

In [1]: `!pip install pyspark`

Requirement already satisfied: pyspark in ./venv/lib/python3.12/site-packages (3.5.5)  
Requirement already satisfied: py4j==0.10.9.7 in ./venv/lib/python3.12/site-packages (from pyspark) (0.10.9.7)

In [2]: `from pyspark.sql import SparkSession`

```
spark = SparkSession.builder \
    .appName("Job Market Analysis 2024") \
    .getOrCreate()
```

Setting default log level to "WARN".  
To adjust logging level use `sc.setLogLevel(newLevel)`. For SparkR, use `setLogLevel(newLevel)`.  
25/04/21 03:48:08 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

In [3]: `df = spark.read.csv("sample_jobs.csv", header=True, inferSchema=True)`

```
df.show(5)
```

```
df.printSchema()
```

```
+-----+-----+-----+-----+-----+-----+-----+
|      JobTitle|  Company|      Location|Salary|  JobType|  Industry|IsAI|
+-----+-----+-----+-----+-----+-----+-----+
|  Data Scientist| TechNova|      Boston|130000|Full-time|Technology|  1|
| Business Analyst|MarketCorp|    Chicago| 85000|Full-time|  Finance|  0|
|    ML Engineer|InnovateAI|San Francisco|150000|Full-time|Technology|  1|
|Software Engineer|  WebWorks|    Seattle|120000|Full-time|Technology|  0|
|   AI Researcher|  DeepMind|    New York|170000|Full-time|Technology|  1|
+-----+-----+-----+-----+-----+-----+-----+
```

only showing top 5 rows

root

```
-- JobTitle: string (nullable = true)
-- Company: string (nullable = true)
-- Location: string (nullable = true)
-- Salary: integer (nullable = true)
-- JobType: string (nullable = true)
-- Industry: string (nullable = true)
-- IsAI: integer (nullable = true)
```

In [4]: `print(f"Rows: {df.count()}, Columns: {len(df.columns)}")`

```
from pyspark.sql.functions import col, isnan, when, count
```

```
df.select([count(when(col(c).isNull() | isnan(c), c)).alias(c) for c in df.columns])
```

Rows: 10, Columns: 7

```
+-----+-----+-----+-----+-----+-----+-----+
|JobTitle|Company|Location|Salary|JobType|Industry|IsAI|
+-----+-----+-----+-----+-----+-----+-----+
|      0|      0|      0|      0|      0|      0|      0|
+-----+-----+-----+-----+-----+-----+-----+
```

```
In [5]: df = df.dropna()
```

```
In [6]: df.show(truncate=False)
```

```
+-----+-----+-----+-----+-----+-----+-----+
|JobTitle      |Company      |Location      |Salary|JobType  |Industry      |IsAI|
+-----+-----+-----+-----+-----+-----+-----+
|Data Scientist |TechNova     |Boston        |130000|Full-time|Technology     |1   |
|Business Analyst|MarketCorp   |Chicago       |85000 |Full-time|Finance        |0   |
|ML Engineer     |InnovateAI   |San Francisco |150000|Full-time|Technology     |1   |
|Software Engineer|WebWorks     |Seattle       |120000|Full-time|Technology     |0   |
|AI Researcher   |DeepMind     |New York      |170000|Full-time|Technology     |1   |
|Marketing Analyst|BrandX       |Austin        |75000 |Full-time|Marketing      |0   |
|Financial Analyst|MoneyMatters |New York      |95000 |Full-time|Finance        |0   |
|Data Engineer   |CloudBase    |Boston        |125000|Full-time|Technology     |1   |
|Project Manager |BuildCo      |Denver        |90000 |Full-time|Construction   |0   |
|AI Product Manager|SmartTech    |San Jose      |145000|Full-time|Technology     |1   |
+-----+-----+-----+-----+-----+-----+-----+
```

```
In [7]: from pyspark.ml.feature import VectorAssembler
```

```
assembler = VectorAssembler(
    inputCols=["Salary", "IsAI"],
    outputCol="features"
)

assembled_data = assembler.transform(df)
assembled_data.select("JobTitle", "features").show(truncate=False)
```

```
+-----+-----+
|JobTitle      |features      |
+-----+-----+
|Data Scientist |[130000.0,1.0]|
|Business Analyst|[85000.0,0.0]|
|ML Engineer     |[150000.0,1.0]|
|Software Engineer|[120000.0,0.0]|
|AI Researcher   |[170000.0,1.0]|
|Marketing Analyst|[75000.0,0.0]|
|Financial Analyst|[95000.0,0.0]|
|Data Engineer   |[125000.0,1.0]|
|Project Manager |[90000.0,0.0]|
|AI Product Manager|[145000.0,1.0]|
+-----+-----+
```

```
In [8]: from pyspark.ml.clustering import KMeans
```

```
kmeans = KMeans(k=2, seed=1, featuresCol="features", predictionCol="cluster")
```

```

model = kmeans.fit(assembled_data)

clustered_data = model.transform(assembled_data)
clustered_data.select("JobTitle", "Salary", "IsAI", "cluster").show(truncate=False)

```

25/04/21 03:48:24 WARN InstanceBuilder: Failed to load implementation from:dev.ludovic.netlib.blas.JNIBLAS

JobTitle	Salary	IsAI	cluster
Data Scientist	130000	1	0
Business Analyst	85000	0	1
ML Engineer	150000	1	0
Software Engineer	120000	0	0
AI Researcher	170000	1	0
Marketing Analyst	75000	0	1
Financial Analyst	95000	0	1
Data Engineer	125000	1	0
Project Manager	90000	0	1
AI Product Manager	145000	1	0

```

In [11]: from pyspark.ml.evaluation import ClusteringEvaluator

clustered_data_for_eval = clustered_data.withColumnRenamed("cluster", "prediction")

evaluator = ClusteringEvaluator(
    featuresCol="features",
    predictionCol="prediction",
    metricName="silhouette",
    distanceMeasure="squaredEuclidean"
)

silhouette = evaluator.evaluate(clustered_data_for_eval)
print(f"Silhouette Score: {silhouette:.3f}")

```

Silhouette Score: 0.806

```

In [12]: from pyspark.ml.feature import StringIndexer, OneHotEncoder, VectorAssembler
from pyspark.ml import Pipeline

indexer = StringIndexer(inputCol="Industry", outputCol="IndustryIndex")
encoder = OneHotEncoder(inputCol="IndustryIndex", outputCol="IndustryVec")
assembler = VectorAssembler(inputCols=["IsAI", "IndustryVec"], outputCol="features")

pipeline = Pipeline(stages=[indexer, encoder, assembler])
pipeline_model = pipeline.fit(df)
transformed_data = pipeline_model.transform(df)

```

```

In [13]: train_data, test_data = transformed_data.randomSplit([0.8, 0.2], seed=42)

```

```

In [14]: from pyspark.ml.regression import LinearRegression

lr = LinearRegression(featuresCol="features", labelCol="Salary")
lr_model = lr.fit(train_data)

```

```
25/04/21 03:49:16 WARN Instrumentation: [e1943a03] regParam is zero, which might cause numerical instability and overfitting.
25/04/21 03:49:17 WARN InstanceBuilder: Failed to load implementation from:dev.ludovic.netlib.lapack.JNILAPACK
25/04/21 03:49:17 WARN Instrumentation: [e1943a03] Cholesky solver failed due to singular covariance matrix. Retrying with Quasi-Newton solver.
```

```
In [15]: predictions = lr_model.transform(test_data)

from pyspark.ml.evaluation import RegressionEvaluator

rmse = RegressionEvaluator(labelCol="Salary", predictionCol="prediction", metricName="rmse")
r2 = RegressionEvaluator(labelCol="Salary", predictionCol="prediction", metricName="r2")

print(f"RMSE: {rmse:.2f}")
print(f"R2: {r2:.2f}")
```

RMSE: 11273.13

R2: 0.85

```
In [16]: lr_model.coefficients
```

```
Out[16]: DenseVector([22500.0008, 45000.0033, 20000.0081, 0.0])
```

```
In [17]: industries = pipeline_model.stages[0].labels

print(f"Intercept: {lr_model.intercept}")
print("Coefficients:")
print(f"IsAI: {lr_model.coefficients[0]}")
for i, name in enumerate(industries):
    print(f"Industry={name}: {lr_model.coefficients[i+1]}")
```

Intercept: 74999.99596926433

Coefficients:

IsAI: 22500.000846654442

Industry=Technology: 45000.00334178167

Industry=Finance: 20000.008119623708

Industry=Construction: 0.0

```
-----
IndexError                                Traceback (most recent call last)
Cell In[17], line 7
      5 print(f"IsAI: {lr_model.coefficients[0]}")
      6 for i, name in enumerate(industries):
----> 7     print(f"Industry={name}: {lr_model.coefficients[i+1]}")

File ~/assignment-02-VidhiSharma2000/myenv/lib/python3.12/site-packages/pyspark/ml/linearalg/__init__.py:469, in DenseVector.__getitem__(self, item)
    468 def __getitem__(self, item: Union[int, slice]) -> Union[np.float64, np.ndarray]:
--> 469     return self.array[item]

IndexError: index 4 is out of bounds for axis 0 with size 4
```

```
In [18]: industries = pipeline_model.stages[0].labels

print(f"Intercept: {lr_model.intercept}")
```

```
print("Coefficients:")
print(f"IsAI: {lr_model.coefficients[0]}")
for i in range(len(lr_model.coefficients) - 1):
    print(f"Industry={industries[i+1]}: {lr_model.coefficients[i+1]}")
```

Intercept: 74999.99596926433  
 Coefficients:  
 IsAI: 22500.000846654442  
 Industry=Finance: 45000.00334178167  
 Industry=Construction: 20000.008119623708  
 Industry=Marketing: 0.0

```
In [20]: indexer = StringIndexer(inputCol="Industry", outputCol="IndustryIndex")
encoder = OneHotEncoder(inputCol="IndustryIndex", outputCol="IndustryVec")
assembler = VectorAssembler(inputCols=["Salary", "IndustryVec"], outputCol="feature")

pipeline = Pipeline(stages=[indexer, encoder, assembler])
pipeline_model = pipeline.fit(df)
transformed_data = pipeline_model.transform(df)
```

```
In [21]: train_data, test_data = transformed_data.randomSplit([0.8, 0.2], seed=42)
```

```
In [22]: from pyspark.ml.classification import LogisticRegression

lr = LogisticRegression(featuresCol="features", labelCol="IsAI")
lr_model = lr.fit(train_data)
```

```
In [23]: predictions = lr_model.transform(test_data)

from pyspark.ml.evaluation import MulticlassClassificationEvaluator

accuracy = MulticlassClassificationEvaluator(labelCol="IsAI", predictionCol="prediction")
f1 = MulticlassClassificationEvaluator(labelCol="IsAI", predictionCol="prediction",

print(f"Accuracy: {accuracy:.2f}")
print(f"F1 Score: {f1:.2f}")
```

Accuracy: 1.00  
 F1 Score: 1.00

```
In [24]: !pip install plotly
```


Collecting plotly  
 Using cached plotly-6.0.1-py3-none-any.whl.metadata (6.7 kB)  
 Collecting narwhals>=1.15.1 (from plotly)  
 Downloading narwhals-1.35.0-py3-none-any.whl.metadata (9.2 kB)  
 Requirement already satisfied: packaging in ./venv/lib/python3.12/site-packages (from narwhals>=1.15.1)  
 Using cached plotly-6.0.1-py3-none-any.whl (14.8 MB)  
 Downloading narwhals-1.35.0-py3-none-any.whl (325 kB)  
 Installing collected packages: narwhals, plotly  
 Successfully installed narwhals-1.35.0 plotly-6.0.1

```
In [25]: import plotly.express as px
import pandas as pd
```

```
pandas_df = df.select("Industry", "Salary").toPandas()

fig = px.box(pandas_df, x="Industry", y="Salary", template="plotly_white", title="S
fig.show()
```

## Salary Distribution by Industry



```
In [26]: pandas_df = df.select("IsAI").toPandas()
pandas_df["IsAI"] = pandas_df["IsAI"].map({1: "AI", 0: "Non-AI"})

fig = px.histogram(pandas_df, x="IsAI", template="plotly_white", title="AI vs Non-A
fig.show()
```


## AI vs Non-AI Job Count



```
In [27]: clustered_df = clustered_data.select("Salary", "IsAI", "cluster").toPandas()

fig = px.scatter(clustered_df, x="Salary", y="IsAI", color="cluster", template="plo
fig.show()
```

## KMeans Job Clustering



## Job Seeker Insights and Recommendations

Based on our analysis of job data from 2024:

- **AI-related jobs** tend to offer significantly higher salaries across all industries, with the average salary in AI roles exceeding non-AI roles by over \$45,000.
- **Industry choice matters** — Technology and Finance roles are high-paying, while roles in Marketing and Construction tend to offer lower compensation.
- **AI classification is highly predictable** from just salary and industry, suggesting a clear separation in job types.
- **Clustering** shows meaningful segmentation of roles, reinforcing that jobs naturally group into high-skill/high-pay and low-skill/low-pay categories.

### Recommendations:

- Job seekers looking to maximize salary potential should **pivot toward AI-focused roles**, especially in the Technology sector.
- Candidates should consider **upskilling with AI and data-related tools** to stand out in the evolving market.



- Non-AI professionals in lower-paying industries should consider **geographic relocation, reskilling, or transitioning industries** to remain competitive.